
MultiRECLAIM Annotator Guide

How to label reclaimed language

Introduction & Golden Rules

Welcome to the MultiRECLAIM annotation thesis project. AI moderation systems frequently penalize LGBTQ+ users for reclaiming their own language, silencing the very voices these systems are supposed to protect. Our goal is to teach the AI the difference between a hostile attack and an empowering reclamation.

Because you self-identify as a member of the LGBTQ+ community or have significant involvement, you possess the sociolinguistic context that AI lacks.

The Golden Rules:

Context is Everything. Your judgment on the speaker's intent matters more than the presence of any single word.

Focus on the Pragmatic Act. Ask yourself: "*Who is speaking, and what are they trying to do with this word?*"

Your Well-being Comes First. You will encounter explicitly homophobic and offensive language. You are never required to annotate a post that causes you distress.

1 The "Skip" Protocol

Because this dataset deals with raw, unfiltered social media data containing slurs, you will encounter highly toxic and upsetting content. You have absolute permission to skip any pair, at any time. Your annotations will still be counted.

Use the skip buttons at the bottom of the task if you cannot or do not want to proceed:

Skip (Distressing). The content is too hostile, or emotionally taxing to evaluate. Use this freely to protect your peace of mind.

Skip (Context). The text is a fragment or a reply that makes absolutely no sense without seeing the previous posts. You cannot safely guess the intent.

Skip (Other). The post is in the wrong language, is pure spam, or is un-annotatable for a reason not listed above.

2 Annotating Text A (The Reclaimed Candidate)

Text A is generated by our AI pipeline as a potential example of reclaimed language. Your job is to verify its quality and intent.

2.1 Step 1: Plausibility Assessment

Does this text sound like authentic social media language?

Highly Plausible. Reads exactly like something a real person would post online.

Plausible, but Awkward. The intent is clear, but the phrasing feels slightly unnatural.

Clearly Synthetic. Reads like an AI manifesto or a dictionary definition.

Critical: Even if you mark a text as synthetic, you must complete Steps 2 and 3. We need to document how the AI's language generation fails.

2.2 Step 2: Pragmatic Action (The Core Task)

Regardless of whether you ultimately think the text is reclaimed, what is the speaker *doing* in this specific context?

Critical: Real-world language is messy. If a text fits multiple categories, you must select the *primary* or *dominant* action driving the post.

Reclamatory / In-Group Uses

Self-Identification. The speaker explicitly claims the word as identity.

Ingroup Solidarity. Casual, affectionate, or unifying peer use.

Ironic / Defiant. Sarcastic or humorous use mocking homophobia.

Non-Reclamatory / Out-Group Uses (Pipeline Traps)

Our AI pipeline sometimes mistakenly accepts the following. Please flag these aggressively:

Victimisation. The speaker is reporting a slur being used against them.

Objectification. Positive sentiment, but from an outsider fetishizing the community.

Metalinguistic. Discussing the word as a concept rather than using it (e.g., “*The word faggot has a long history*”).

Anger / Venting. An insider expressing frustration about the word or the community.

Internalised Negative. An insider using the word against themselves with genuine shame.

Hostile Attack. Weaponized, aggressive use intended to harm.

Unclear

If the text could genuinely be read multiple ways or is too fragmented, select **Unclear** and specify the reason.

Show

You do not need to memorize these definitions! If you ever need a quick refresher while annotating, click the show button at the bottom of the label list. This will expand a quick-reference glossary directly inside your task view.

2.3 Step 3: Final Verdict & Logical Mapping

Based on your Step 2 analysis, is Text A a valid example of reclaimed usage?

Note: Your verdict here should logically follow Step 2. If you selected any of the *Pipeline Traps* (Victimisation, Hostile Attack, etc.), your Final Verdict should almost always be *Not Reclaimed*.

Reclaimed (select *High*, *Medium*, or *Low* confidence). Use this if you selected **Self-Identification**, **Ingroup Solidarity**, or **Ironic/Defiant** in Step 2.

Not Reclaimed (select *High*, *Medium*, or *Low* confidence). Use this if you selected **Victimisation**, **Objectification**, **Metalinguistic**, **Hostile Attack**, or **Internalised Negative** in Step 2.

Ambiguous (select *High*, *Medium*, or *Low* confidence).

3 Annotating Text B (The Generated Counterpart)

Text B is generated by the AI to act as a contrasting counterpart to Text A. The goal is to see if the AI successfully shifted the pragmatic intent (from reclaimed to non-reclaimed).

3.1 Step 4a: Counterpart Plausibility

Evaluate Text B’s authenticity exactly as you did for Text A: **Highly Plausible**, **Plausible/Awkward**, or **Clearly Synthetic**.

3.2 Step 4b: Contrastive Utility Validation

Note: Do not worry about counting words or checking exact sentence structures. The automated pipeline already handles structural similarity. We only need your human judgment on the intent.

Does Text B successfully shift the stance to *Non-Reclaimed*?

Strong Contrast. The intent successfully flipped to a non-reclaimed stance.

Valid Loose. The intent successfully flipped to a non-reclaimed stance, but the AI wrote a completely different sentence to get there.

Failed Contrast. The AI failed to shift the intent (e.g., Text B is still a reclaimed use, or the text is complete nonsense).



Scan to open registration

Questions?

Contact us at contact@multireclaim.study

Remember: Your well-being comes first.

Skipping is always allowed.